SHORT · COPY INTO

COPY INTO rerun loads 0 rows

COPY INTO is a SQL command that loads files from a folder into an existing Delta table. It is idempotent: Databricks keeps a load history for each target table, and files already loaded into that table are skipped on the next run. So a rerun with no new files inserts 0 rows. COPY_OPTIONS ('force' = 'true') turns idempotency off and loads every file again.

Key points

- Rerun with no new files: num_inserted_rows 0, total stays 3. No duplicates.
- orders_day2.csv lands: only its 2 rows load, total 5. Each file loads exactly once.
- Edited file with the same name: still skipped.
- New target table: it has its own load history, so all files load into it.
- force = true goes in COPY_OPTIONS: every file loads again. After day 2: 5 inserted, total 10, duplicates.

Tiny example

A03 example, fictional orders in /Volumes/workspace/bt_copyinto_xxxx/landing/orders. Run 1: 3 rows. Run 2: 0 rows. Run 3: 2 rows, total 5.

```
COPY INTO orders_copy
FROM '/Volumes/workspace/bt_copyinto_xxxx/landing/orders'
FILEFORMAT = CSV
FORMAT_OPTIONS ('header' = 'true');
-- rerun, no new files: 0 rows
-- COPY_OPTIONS ('force' = 'true'): every file again
```

Misconception check

COPY INTO does not append every file on each run. The trap answer is force = true: it turns idempotency off, so every file loads again and the table gets duplicates.

Memory trick

COPY INTO remembers. force forgets.

Remember

No new files = 0 rows. force = true = every file again, duplicates included.

Expected results from the A03 lesson notebook (Databricks Free Edition, serverless notebook, Unity Catalog volume). Full lesson: lessons/associate-certification/a03-copy-into.